


# Reinforcement Learning workshop

Florian Goebels

October 3ed 2019



# Agenda

- Day 1: RL Basics (theory)
- Day 2: Q-learning (coding)
- Day 3: Policy Gradient methods (coding)

# Today's Agenda

- Day 3: Policy Gradients
  - Background
  - Deep deterministic deep policy gradients

# Q-learning

Q-learning (off-policy TD control) for estimating  $\pi \approx \pi_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

        Take action  $A$ , observe  $R, S'$

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

    until  $S$  is terminal

# TD $V(S)$ estimate

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$



Better estimate of  $V(S)$

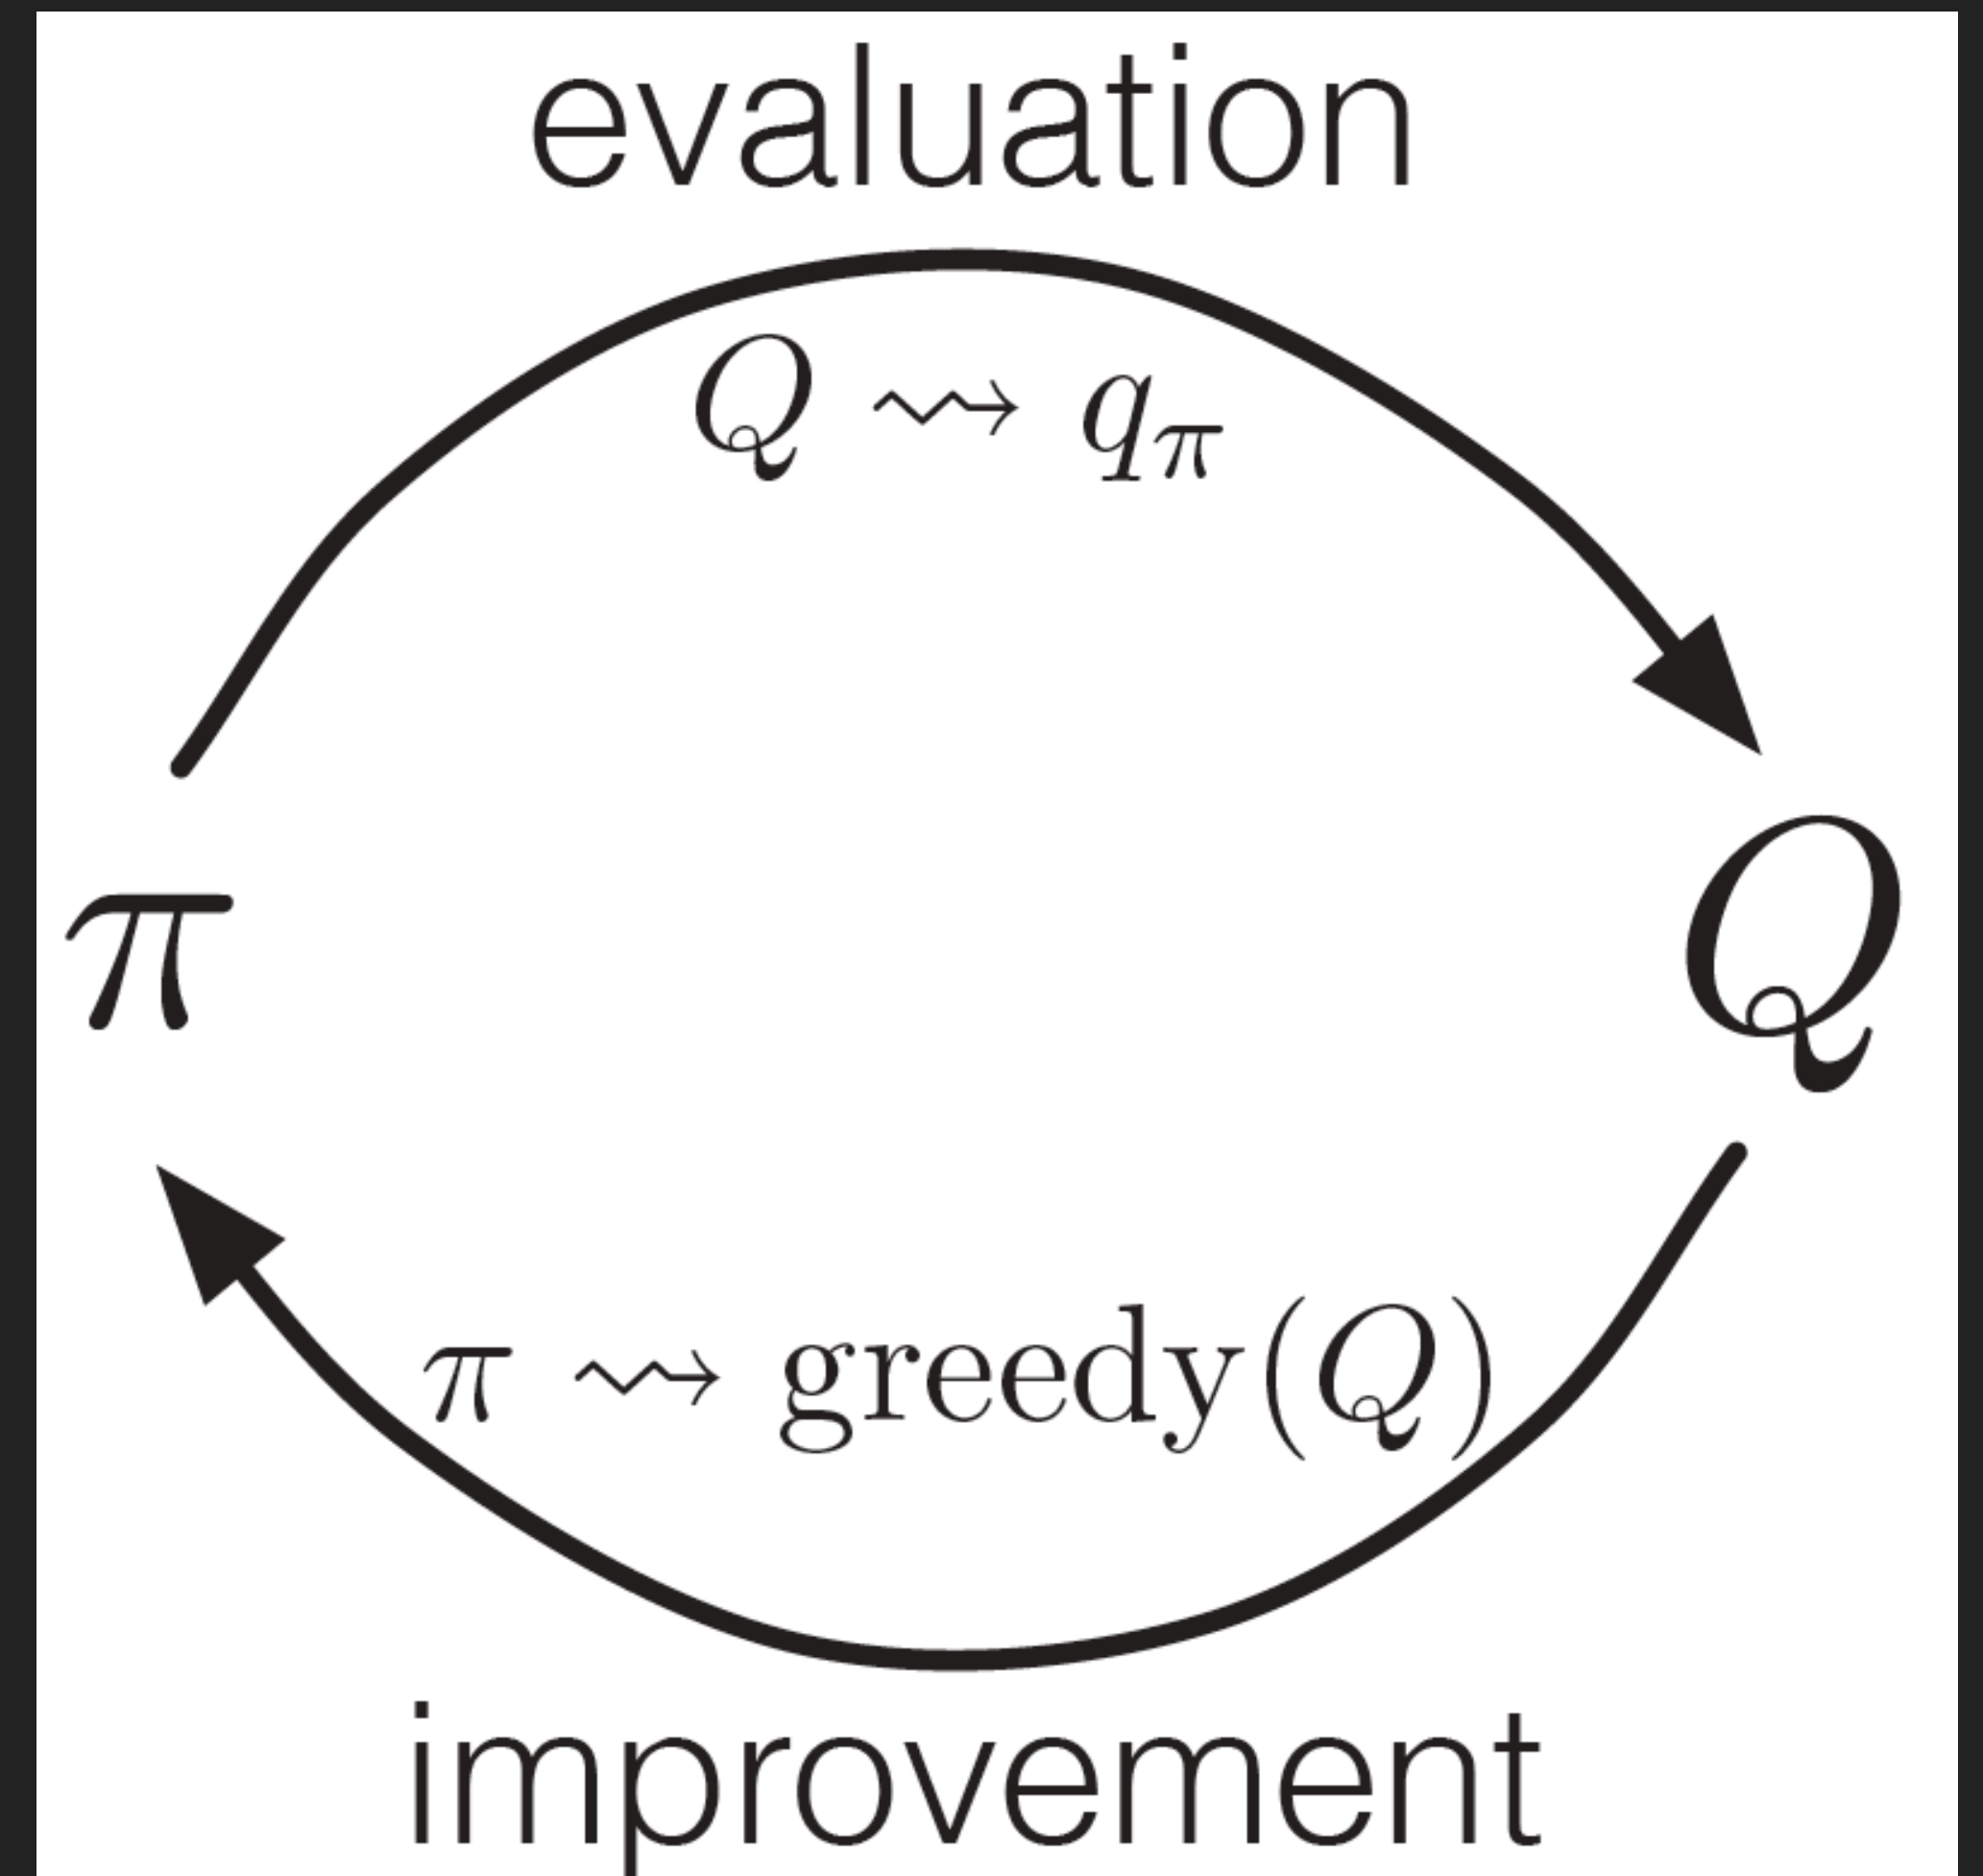
$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

TD ERROR

Current estimate of  $V(S)$

# Generalized Policy iteration

- Almost all RL algorithm are GPI
- Maintain both an approximate value function and an approximate policy
- Iteratively improve policy with respect to value function, and value function always drives to the value function of the current policy
- Overall process converges to optimal policy and optimal value function



# Deep Q-Learning

- In Q learning we iteratively update  $Q(A,S)$  via policy iteration
- Neural networks are general function approximators  $NN(S) \rightarrow V(A)$
- Cost function:

Better estimate of  $V(S)$

$$\left( \text{NN}(S_t, a) - r + \gamma \max_a \text{NN}(S_{t+1}, a) \right)^2$$

NN estimation of  $V(S)$

# Policy Gradient


Weight matrix  $\theta \in R^{d \times |A|}$       State features  $x(s, a) \in R^{1 \times d}$

Hidden  $h(s, a, \theta) = \theta \times x(s, a)$

Softmax activation  $\pi(a \mid s, \theta) = \frac{e^{h(s, a, \theta)}}{\sum e^{h(s, a, \theta)}}$



# Neural network propagation



**Forward:**

$$X * W = H$$

$$\sigma(H) = \text{softmax}(H) = Y$$

**Backward:**

$$\text{Cost function : } C(y, y') = (y - y')^2$$

$$\frac{\partial C(y, y')}{\partial W} = X^T * (y' - y) * \sigma'(H)$$

# Regression towards optimal policy

$$SW = \pi$$

$$SW = [0.2, 0.8]$$

$$\pi^* = [0.4, 0.6]$$

$$C([0.2, 0.8], [0.4, 0.6]) = 0.28$$

**We don't know optimal policy**

# Policy gradient

Cost function

$$J(\theta)$$

Gradient ascent

$$\theta_{t+1} = \theta_t + \alpha \nabla_{\theta} J(\theta)$$

Policy gradient

$$\nabla_{\theta} J(\theta) = G \nabla \ln \pi(a | s, \theta)$$

# REINFORCE

REINFORCE: Monte-Carlo Policy-Gradient Control (episodic) for  $\pi_*$

Input: a differentiable policy parameterization  $\pi(a|s, \theta)$

Algorithm parameter: step size  $\alpha > 0$

Initialize policy parameter  $\theta \in \mathbb{R}^{d'}$  (e.g., to  $\mathbf{0}$ )

Loop forever (for each episode):

    Generate an episode  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$ , following  $\pi(\cdot|\cdot, \theta)$

    Loop for each step of the episode  $t = 0, 1, \dots, T - 1$ :

$$G \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k \quad (G_t)$$

$$\theta \leftarrow \theta + \alpha \gamma^t G \nabla \ln \pi(A_t|S_t, \theta)$$

# Policy gradient

$$\nabla_{\theta} J(\theta) = G \nabla \ln \pi(a | s, \theta) = \frac{\nabla \pi(a, s, \theta)}{\pi(a, s, \theta)}$$

$$\nabla \pi(a, s, \theta) = \frac{\partial \sigma(XW)}{\partial W} =$$

$$\frac{\partial \sigma(XW)}{\partial XW} * \frac{\partial XW}{W} = X^T * \sigma'(XW)$$

# Deep deterministic policy gradient

# Actor critic

Actor: Policy function controls our agent

$$\pi(A, A, \theta)$$

Critic: Value function measures how good  
actions are

$$q(s, a, w)$$



# DDPG

---

**Algorithm 1** DDPG algorithm

---

Randomly initialize critic network  $Q(s, a|\theta^Q)$  and actor  $\mu(s|\theta^\mu)$  with weights  $\theta^Q$  and  $\theta^\mu$ .

Initialize target network  $Q'$  and  $\mu'$  with weights  $\theta^{Q'} \leftarrow \theta^Q, \theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer  $R$

**for** episode = 1, M **do**

    Initialize a random process  $\mathcal{N}$  for action exploration

    Receive initial observation state  $s_1$

**for** t = 1, T **do**

        Select action  $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$  according to the current policy and exploration noise

        Execute action  $a_t$  and observe reward  $r_t$  and observe new state  $s_{t+1}$

        Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $R$

        Sample a random minibatch of  $N$  transitions  $(s_i, a_i, r_i, s_{i+1})$  from  $R$

        Set  $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$

        Update critic by minimizing the loss:  $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

        Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

    Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

**end for**

**end for**

---

# DDPG

---

**Algorithm 1** DDPG algorithm

---

Randomly initialize critic network  $Q(s, a|\theta^Q)$  and actor  $\mu(s|\theta^\mu)$  with weights  $\theta^Q$  and  $\theta^\mu$ .  
Initialize target network  $Q'$  and  $\mu'$  with weights  $\theta^{Q'} \leftarrow \theta^Q, \theta^{\mu'} \leftarrow \theta^\mu$

- Initialize:
  - Actor: input state - output policy
  - Critic: input state + policy - output q-value

# DDPG

```
for episode = 1, M do  
  Initialize a random process  $\mathcal{N}$  for action exploration  
  Receive initial observation state  $s_1$   
  for t = 1, T do  
    Select action  $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$  according to the current policy and exploration noise  
    Execute action  $a_t$  and observe reward  $r_t$  and observe new state  $s_{t+1}$   
    Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $R$ 
```

- Play game and use actor to create policy and store transitions in replay memory

# DDPG

Sample a random minibatch of  $N$  transitions  $(s_i, a_i, r_i, s_{i+1})$  from  $R$

Set  $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$

Update critic by minimizing the loss:  $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

- Sample mini batch
- predict policy from state with  $Q'$  and predict value of state with  $\mu'$
- Calculate loss between  $y$  and model prediction of historic state, action pair

# DDPG

Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a | \theta^Q) |_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s | \theta^\mu) |_{s_i}$$

- Calculate policy gradient for actor

# DDPG

Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^{\mu} + (1 - \tau) \theta^{\mu'}$$

- Update networks



# License

All course material- including this slide deck, hands-on workbooks, and any other guided exercise for example assignments- are part of the workshop "RL workshop" and owned by Aggregate Intellect Inc. (<https://ai.science>), and released under "Creative Commons Attribution-NonCommercial-ShareAlike CC BY-NC-SA" license.

This material can be altered and distributed for non-commercial use with reference to Aggregate Intellect Inc. as the original owner, and any material generated from it must be released under similar terms.

For more details see: <https://creativecommons.org/licenses/by-nc-sa/4.0/>